

Linear Models in R (M1–MIDO)

Lab Session 3 – Solutions

Henri PANJO

Table of contents

Dataset Overview: <i>data_pokemon.csv</i>	3
Setup	4
Question 1. Loading dataset	5
Solutions	5
Question 2: Exploring categorical variables	6
Solutions	6
Question 3: Data management, variable creation	11
Solutions	11
Question 4: Box plot	15
Solutions	15
Question 5: Attack mean over group variables	17
Solutions	17
Question 6: Dummy variables creation	20
Solutions	20
Question 7: Using dummy variables in OLS regression	21
Solutions	21
Question 8: Testing equality of coefficients	27
Solutions	27
Question 9: Predicting mean attack for all category combinations	29
Solutions	29
Question 10. Residual diagnostics	32
Solutions	32
Session Info	36

Dataset Overview: `data_pokemon.csv`

This dataset is adapted from a popular Kaggle Pokéémon dataset.

Even if you are not familiar with Pokéémon, the data is straightforward:

it combines numeric statistics with categorical attributes, making it well-suited for applying Ordinary Least Squares (OLS) in R.

What it contains

- Unique identifiers and names for each Pokéémon
- Battle statistics (health, attack, defense, special attack, special defense, speed)
- Categorical features (primary/secondary type, generation, legendary flag)

Fields (Codebook)

- `id`: Unique Pokéémon ID
- `name`: Pokéémon name
- `type_1`: Primary type (e.g., Water, Fire)
- `type_2`: Secondary type (optional)
- `hp`: Hit points (overall health)
- `attack`: Physical attack strength (we will use this as y in most regressions)
- `defense`: Physical defense strength
- `sp_attack`: Special (non-physical) attack strength
- `sp_defense`: Special defense strength
- `speed`: Speed / turn order
- `generation`: Game generation label
- `legendary`: Indicator for legendary status (TRUE/FALSE)

Note on notation

- We treat `attack` as the outcome variable Y .
- Predictor variables (e.g., `defense`, `speed`) will be denoted as x_1, x_2, \dots .
- Factors like `type_1` or `legendary` will be included as categorical predictors.

Setup

To keep numbers readable and reproducible, we set display options:

```
options(scipen = 999, digits = 5)
```

We also load the packages used during this session.

⚠ Warning

Don't worry if you don't know them all — we'll introduce functions as we need them. Some provide regression tools, others are for data visualization or diagnostics.

```
library(broom)
library(performance)
library(parameters)
library(datawizard)
library(see)
library(effectsize)
library(insight)
library(correlation)
library(modelbased)
library(glue)
library(scales)
library(GGally)
library(ggpubr)
library(car)
library(lmtest)
library(multcomp)
library(rstatix)
library(matrixTests)
library(ggfortify)
library(qqplotr)
library(patchwork)
library(gtsummary)
library(kableExtra)
library(openxlsx)
library(janitor)
library(collapse)
library(tidyverse)

source("helper_functions3.R")
```

Question 1. Loading dataset

Import the dataset `data_pokemon.csv` with `read_csv()` and save it in an object called `pok`.

Using `select()`, keep only the variables `id`, `name`, `attack`, `speed`, `defense`, `hp`, `sp_attack`, and `sp_def`.

Display the first 10 rows of `pok` using `head()` or `slice()`.

Solutions

- Loading `data_pokemon.csv`

```
pok <- read_csv("data_pokemon.csv", show_col_types = FALSE)
```

- `head()` on `pok`

```
head(pok, n = 10)
```

```
# A tibble: 10 x 12
  id name   type_1 type_2   hp attack defense sp_attack sp_def speed generation legendary
  <dbl> <chr>  <chr>  <chr>  <dbl> <dbl>   <dbl>    <dbl>   <dbl>   <dbl>   <dbl> <chr>
1     1 Bulbasaur Grass  Poison    45     49     49      65     65     45      1 No
2     2 Ivysaur   Grass  Poison    60     62     63      80     80     60      1 No
3     3 Venusaur  Grass  Poison    80     82     83     100    100     80      1 No
4     4 Mega Ven~ Grass  Poison    80    100    123     122    120     80      1 No
5     5 Charmand~ Fire   None     39     52     43      60     50     65      1 No
6     6 Charmele~ Fire   None     58     64     58      80     65     80      1 No
7     7 Charizard  Fire  Flying    78     84     78     109     85    100      1 No
8     8 Mega Cha~ Fire  Dragon    78    130    111     130     85    100      1 No
9     9 Mega Cha~ Fire  Flying    78    104     78     159    115    100      1 No
10    10 Squirtle  Water  None     44     48     65      50     64     43      1 No
```

Question 2: Exploring categorical variables

The Pokémon dataset contains 4 categorical variables

- type_1 : the primary type (always present)
- type_2 : the secondary type (may be missing)
- legendary: if the pokémon is legendary or not.
- generation: the pokémon generation.

1. Create frequency tables for type_1, type_2, legendary and generation. Display both the counts and the relative proportions.
2. Produce bar plots for the distributions of type_1 and type_2.
 - Make one bar plot for type_1 and one for type_2.
 - Ensure that categories on the x-axis are readable (e.g., rotate labels if necessary).

Solutions

- Frequency table for type_1 with `count()` from `{dplyr}`

```
table_type1_prop <- mutate(pok, type_1 = fct_infreq(type_1)) |>
  count(type_1) |>
  mutate(pct = 100 * n / sum(n))
```

```
table_type1_prop
```

	type_1	n	pct
1	Water	112	14
2	Normal	98	12.2
3	Grass	70	8.75
4	Bug	69	8.62
5	Psychic	57	7.12
6	Fire	52	6.5
7	Electric	44	5.5
8	Rock	44	5.5
9	Dragon	32	4
10	Ghost	32	4
11	Ground	32	4
12	Dark	31	3.88
13	Poison	28	3.5
14	Fighting	27	3.38
15	Steel	27	3.38
16	Ice	24	3
17	Fairy	17	2.12
18	Flying	4	0.5

- Frequency table for type_2 with `tabyl()` from `{janitor}`

```
table_type2_prop <- tabyl(pok, type_2) |>
  adorn_pct_formatting(digits = 2) |>
  arrange(desc(n))
```

`table_type2_prop`

type_2	n	percent
None	386	48.25%
Flying	97	12.12%
Ground	35	4.38%
Poison	34	4.25%
Psychic	33	4.12%
Fighting	26	3.25%
Grass	25	3.12%
Fairy	23	2.88%
Steel	22	2.75%
Dark	20	2.50%
Dragon	18	2.25%
Ghost	14	1.75%
Ice	14	1.75%
Rock	14	1.75%
Water	14	1.75%
Fire	12	1.50%
Electric	6	0.75%
Normal	4	0.50%
Bug	3	0.38%

- Frequency table for legendary and generation with `tabyl()`

```
tabyl(pok, legendary) |>
  adorn_pct_formatting(digits = 2)
```

legendary	n	percent
No	735	91.88%
Yes	65	8.12%

```
tabyl(pok, generation) |>
  adorn_pct_formatting(digits = 2)
```

generation	n	percent
1	166	20.75%
2	106	13.25%
3	160	20.00%
4	121	15.12%
5	165	20.62%
6	82	10.25%

- Frequency tables for type_1 and type_2 with `tbl_summary()` from `{gtsummary}`

```
select(pok, type_1) |>
  mutate(type_1 = fct_infreq(type_1)) |>
  tbl_summary(label = list(type_1 = "Pokemon primary type")) |>
  modify_header(all_stat_cols() ~ "**{level} (n={n})**") |>
  bold_labels() |>
  remove_footnote_header()

select(pok, type_2) |>
  mutate(type_2 = fct_infreq(type_2)) |>
  tbl_summary(label = list(type_2 = "Pokemon secondary type")) |>
  modify_header(all_stat_cols() ~ "**{level} (n={n})**") |>
  bold_labels() |>
  remove_footnote_header()
```

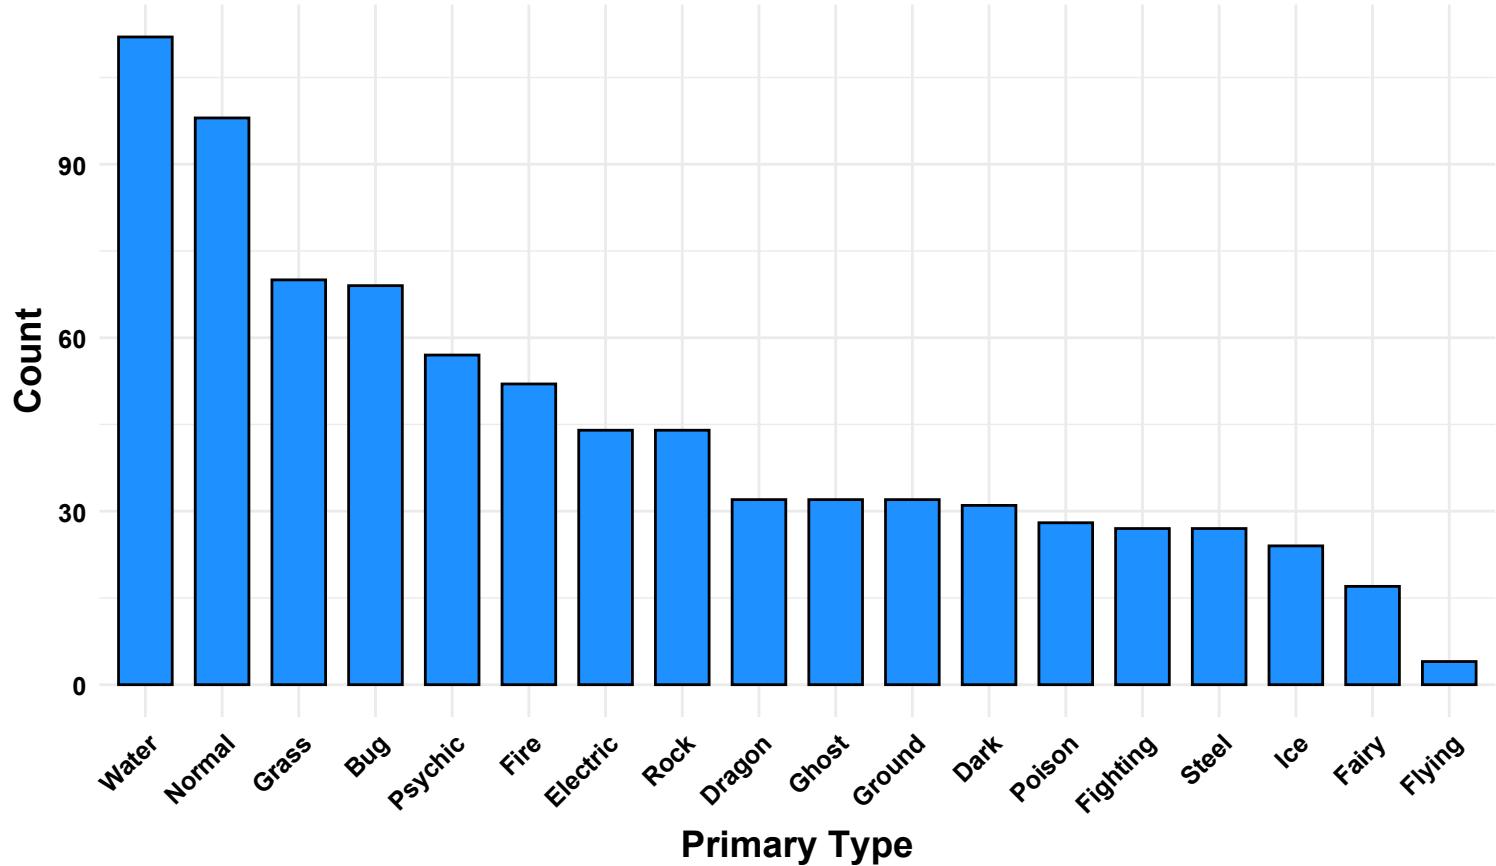
Characteristic	Overall (n=800)
Pokemon primary type	
Water	112 (14%)
Normal	98 (12%)
Grass	70 (8.8%)
Bug	69 (8.6%)
Psychic	57 (7.1%)
Fire	52 (6.5%)
Electric	44 (5.5%)
Rock	44 (5.5%)
Dragon	32 (4.0%)
Ghost	32 (4.0%)
Ground	32 (4.0%)
Dark	31 (3.9%)
Poison	28 (3.5%)
Fighting	27 (3.4%)
Steel	27 (3.4%)
Ice	24 (3.0%)
Fairy	17 (2.1%)
Flying	4 (0.5%)

Characteristic	Overall (n=800)
Pokemon secondary type	
None	386 (48%)
Flying	97 (12%)
Ground	35 (4.4%)
Poison	34 (4.3%)
Psychic	33 (4.1%)
Fighting	26 (3.3%)
Grass	25 (3.1%)
Fairy	23 (2.9%)
Steel	22 (2.8%)
Dark	20 (2.5%)
Dragon	18 (2.3%)
Ghost	14 (1.8%)
Ice	14 (1.8%)
Rock	14 (1.8%)
Water	14 (1.8%)
Fire	12 (1.5%)
Electric	6 (0.8%)
Normal	4 (0.5%)
Bug	3 (0.4%)

- Bar plots for the distribution of type_1 with `geom_bar()`


```
pok |>
  mutate(type_1 = fct_infreq(type_1)) |>      # order by descending frequency
  ggplot(aes(x = type_1)) +
  geom_bar(fill = "dodgerblue", color = "black", width = 0.7) +
  labs(
    title = "Distribution of Primary Pokémon Types",
    x = "Primary Type", y = "Count"
  ) +
  theme_minimal() +
  theme(axis.text.x = element_text(angle = 45, hjust = 1)) +
  labs_pubr()
```

Distribution of Primary Pokémon Types



- Bar plots for the distribution of type_1 with `geom_col()`

```
pok |>
  mutate(type_2 = fct_infreq(type_2)) |> # order by descending frequency
  count(type_2) |> # get the count
  ggplot(aes(y = type_2, x = n, label = n)) +
  geom_col(fill = "dodgerblue", color = "black", width = 0.7) +
  geom_text(hjust = -0.2) +
  labs(
    title = "Distribution of Primary Pok  mon Types",
    x = "Count", y = "Secondary Type"
  ) +
  theme_minimal() +
  labs_pubr()
```



Question 3: Data management, variable creation

The variable (type_1), has 18 levels, which can be too many to include directly in an OLS regression. To simplify the analysis, we want to group these 18 types into 3 broader, meaningful groups:

- Physical/Material: Bug, Fighting, Ground, Rock, Steel, Normal
- Elemental/Environmental: Fire, Water, Grass, Electric, Ice, Flying, Poison
- Mystical/Supernatural: Psychic, Ghost, Dragon, Fairy, Dark

1. In the pok dataframe, create a new factor variable called type_group3 that assigns each Pokémon to one of the 3 groups above based on its primary type (type_1).
2. Create a new binary variable called has_secondary_type defined as follows:
 - “Yes” if the Pokémon has a secondary type (type_2 is not “None”)
 - “No” if the Pokémon does not have a secondary type (type_2 is “None”)
3. Transform the variables legendary and generation into factors.
4. Verify that the new variables are well created

Solutions

Creation of the variable type_group3 and has_secondary_type

- Define the 3-level grouping map

```
type_map3 <- list(  
  physical_material = c("Bug", "Fighting", "Ground", "Rock", "Steel", "Normal"),  
  elemental_env = c("Fire", "Water", "Grass", "Electric", "Ice", "Flying", "Poison"),  
  mystical_supernatural = c("Psychic", "Ghost", "Dragon", "Fairy", "Dark"))
```

- Create variables with `mutate()` from `{dplyr}`

```
pok <- mutate(  
  pok,  
  type_group3 = case_when(  
    type_1 %in% type_map3$physical_material ~ "Physical/Material",  
    type_1 %in% type_map3$elemental_env ~ "Elemental/Environmental",  
    type_1 %in% type_map3$mystical_supernatural ~ "Mystical/Supernatural",  
    .default = NA_character_)  
,  
  has_secondary_type = ifelse(type_2 == "None", 0, 1) |> factor(labels = c("No", "Yes"))  
) |>  
  mutate(type_group3 = fct_infreq(type_group3), legendary = factor(legendary)) |>  
  mutate(generation = factor(generation, labels = paste0("G", 1:6))) |>  
  relabel(  
    type_group3 = "Primary Type", has_secondary_type = "Has a secondary type",  
    legendary = "Legendary", generation = "Pokemon generation")  
)
```

- type_group3: variable checking

```
tabyl(pok, type_group3) |> as_tibble()
```

```
# A tibble: 3 x 3
  type_group3          n percent
  <fct>              <int>   <dbl>
1 Elemental/Environmental 334    0.418
2 Physical/Material      297    0.371
3 Mystical/Supernatural 169    0.211
```

```
tabyl(pok, type_1, type_group3) |> as_tibble()
```

```
# A tibble: 18 x 4
  type_1    `Elemental/Environmental` `Physical/Material` `Mystical/Supernatural`
  <chr>        <dbl>                  <dbl>                  <dbl>
1 Bug            0                    69                    0
2 Dark           0                    0                     31
3 Dragon          0                    0                     32
4 Electric         44                   0                     0
5 Fairy            0                    0                     17
6 Fighting          0                   27                     0
7 Fire             52                   0                     0
8 Flying            4                    0                     0
9 Ghost            0                    0                     32
10 Grass            70                   0                     0
11 Ground           0                   32                     0
12 Ice              24                   0                     0
13 Normal            0                   98                     0
14 Poison            28                   0                     0
15 Psychic           0                   0                     57
16 Rock              0                   44                     0
17 Steel             0                   27                     0
18 Water            112                  0                     0
```

- has_secondary_type: variable checking

```
tabyl(pok, has_secondary_type) |> as_tibble()
```

```
# A tibble: 2 x 3
  has_secondary_type      n percent
  <fct>              <int>   <dbl>
1 No                  386    0.482
2 Yes                 414    0.518
```

```
tabyl(pok, type_2, has_secondary_type) |> as_tibble()
```

```
# A tibble: 19 x 3
  type_2      No    Yes
  <chr>     <dbl> <dbl>
1 Bug          0     3
2 Dark         0    20
3 Dragon        0    18
4 Electric       0     6
5 Fairy         0    23
6 Fighting       0    26
7 Fire          0    12
8 Flying         0    97
9 Ghost          0    14
10 Grass         0    25
11 Ground        0    35
12 Ice           0    14
13 None          386   0
14 Normal         0     4
15 Poison         0    34
16 Psychic        0    33
17 Rock           0    14
18 Steel          0    22
19 Water          0    14
```

- Cross tabulation between type_group3 and has_secondary_type

```
tabyl(pok, type_group3, has_secondary_type) |>
  adorn_totals(where = c("row", "col")) |>
  adorn_percentages(denominator = "row") |>
  adorn_pct_formatting(digits = 1) |>
  adorn_ns(position = "front") |>
  as_tibble() |>
  mutate(across(everything(), str_squish))
```

```
# A tibble: 4 x 4
  type_group3      No       Yes      Total
  <chr>     <chr>     <chr>     <chr>
1 Elemental/Environmental 177 (53.0%) 157 (47.0%) 334 (100.0%)
2 Physical/Material      125 (42.1%) 172 (57.9%) 297 (100.0%)
3 Mystical/Supernatural   84 (49.7%)  85 (50.3%) 169 (100.0%)
4 Total                  386 (48.2%) 414 (51.7%) 800 (100.0%)
```

- legendary and generation : variable checking

```
count(pok, legendary)
```

```
# A tibble: 2 x 2
  legendary     n
  <fct>    <int>
1 No        735
2 Yes       65
```

```
count(pok, generation)
```

```
# A tibble: 6 x 2
  generation     n
  <fct>    <int>
1 G1        166
2 G2        106
3 G3        160
4 G4        121
5 G5        165
6 G6         82
```

Question 4: Box plot

We now want to explore how the Pokémon attack distribution varies across several categorical variables in the dataset. Boxplots are useful for comparing the distribution of a numeric variable across groups.


Using the Pokémon dataset, create four separate boxplots where the response variable is attack, and the grouping variables are `type_group3`, `has_secondary_type`, `legendary`, `generation`.

Produce the four boxplots either separately or arranged in a multi-panel layout (your choice).

Solutions

- Boxplots of attack by `type_group3`

```
ggplot(pok, aes(x = type_group3, y = attack)) +  
  geom_boxplot(linewidth = 0.25, median.linewidth = 0.75) +  
  labs(y = "Pokemon power attack", x = vlabels(pok$type_group3)) +  
  theme_bw(base_size = 14) +  
  labs_pubr()
```




- Boxplots in a multi-panel layout

```
varcats <- c("type_group3", "has_secondary_type", "legendary", "generation")

boxplots <- select(pok, id, attack, all_of(varcats)) |>
  pivot_longer(all_of(varcats), names_to = "var") |>
  mutate(var = factor(var, levels = varcats, labels = vlabels(pok[, varcats]))) |>
  mutate(value = fct_relabel(value, \((x) str_trunc(x, 12)))) |>
  ggplot(aes(x = attack, y = value)) +
  geom_boxplot(linewidth = 0.25, median.linewidth = 0.75) +
  facet_wrap(vars(var), scales = "free_y") +
  labs(x = "Pokemon power attack", y = NULL) +
  theme_bw(base_size = 14) +
  labs_pubr() +
  theme(
    strip.text = element_text(size = 11, face = "bold"),
    axis.text.y = element_text(size = 8, face = "bold")
  )

```

boxplots



Question 5: Attack mean over group variables

You have explored the distribution of the variable `attack` using boxplots. We now want to summarize these differences numerically by computing the mean Attack score for several grouping variables.

Using the updated Pokémon dataset, compute the mean value of `attack` for each level of the following categorical variables:

- `type_group3` (3-level grouped primary type)
- `has_secondary_type` ("No" / "Yes")
- `legendary` ("No" / "Yes")
- `generation` ("G1" to "G6")

For each variable, produce a summary table showing at least:

- the group name
- the mean of `attack`
- the standard deviation of `attack`
- the number of observations in each group

Hint: Use `mean_by_group()` from `helper_functions3.R`

Solutions

- We use `mean_by_group()`

```
mean_by_group(data = pok, x = "attack", by = "type_group3")
```

```
# A tibble: 4 x 3
  Variable           N `Mean (SD)`
  <chr>              <int> <chr>
1 Primary Type       NA <NA>
2 Elemental/Environmental 334 74.9 (26.8)
3 Physical/Material   297 82.0 (34.4)
4 Mystical/Supernatural 169 81.7 (38.0)
```

```
mean_by_group(pok, "attack", "has_secondary_type")
```

```
# A tibble: 3 x 3
  Variable           N `Mean (SD)`
  <chr>              <int> <chr>
1 Has a secondary type NA <NA>
2 No                  386 74.5 (30.5)
3 Yes                 414 83.2 (33.7)
```

```
mean_by_group(pok, "attack", "legendary")
```

```
# A tibble: 3 x 3
  Variable      N `Mean (SD)`
  <chr>        <int> <chr>
1 Legendary     NA <NA>
2 No            735 75.7 (30.5)
3 Yes           65 116.7 (30.3)
```

```
mean_by_group(pok, "attack", "generation", digits = 2)
```

```
# A tibble: 7 x 3
  Variable      N `Mean (SD)`
  <chr>        <int> <chr>
1 Pokemon generation NA <NA>
2 G1            166 76.64 (30.74)
3 G2            106 72.03 (32.71)
4 G3            160 81.63 (36.59)
5 G4            121 82.87 (32.78)
6 G5            165 82.07 (30.37)
7 G6            82 75.80 (29.18)
```

- We can use `mean_by_group()` on different groups with `split()` from Base R and `map()` from `{purrr}`

```
split(pok, ~legendary) |>
  map(\(d) mean_by_group(d, x = "attack", by = "type_group3", digits = 2)) |>
  list_rbind(names_to = "Legendary Pokemon")
```

```
# A tibble: 8 x 4
`Legendary Pokemon` Variable      N `Mean (SD)`
<chr>        <chr>        <int> <chr>
1 No          Primary Type    NA <NA>
2 No          Elemental/Environmental 314 73.25 (26.25)
3 No          Physical/Material   283 79.89 (33.20)
4 No          Mystical/Supernatural 138 72.51 (32.82)
5 Yes         Primary Type    NA <NA>
6 Yes         Elemental/Environmental 20 101.40 (21.08)
7 Yes         Physical/Material   14 125.29 (30.59)
8 Yes         Mystical/Supernatural 31 122.65 (32.42)
```

- All in one table using `map()` from `{purrr}`

```
tabmeans <- map(varcats, \by) mean_by_group(pok, "attack", by)) |>
  list_rbind() |>
  mutate(Variable = ifelse(is.na(N), Variable, paste0("\u2000", Variable)))

tabmeans |>
  kable(align = "lcc", padding = 2) |>
  row_spec(c(1, 5, 8, 11), bold = TRUE, color = "white", background = "black")
```

Variable	N	Mean (SD)
Primary Type		
Elemental/Environmental	334	74.9 (26.8)
Physical/Material	297	82.0 (34.4)
Mystical/Supernatural	169	81.7 (38.0)
Has a secondary type		
No	386	74.5 (30.5)
Yes	414	83.2 (33.7)
Legendary		
No	735	75.7 (30.5)
Yes	65	116.7 (30.3)
Pokemon generation		
G1	166	76.6 (30.7)
G2	106	72.0 (32.7)
G3	160	81.6 (36.6)
G4	121	82.9 (32.8)
G5	165	82.1 (30.4)
G6	82	75.8 (29.2)

- To export the table in Excel format, you can use `write.xlsx()` from `{openxlsx}`

```
write.xlsx(tabmeans, "tabmeans.xlsx")
```

Question 6: Dummy variables creation

In the Pokémon dataset create these variables

$$\text{legend1} = \begin{cases} 1 & \text{if } \text{legendary} = \text{"Yes"} \\ 0 & \text{otherwise} \end{cases} \quad \text{legend0} = \begin{cases} 1 & \text{if } \text{legendary} = \text{"No"} \\ 0 & \text{otherwise} \end{cases}$$

$$\text{typeg1} = \begin{cases} 1 & \text{if } \text{type_group3} = \text{"Elemental/Environmental"} \\ 0 & \text{otherwise} \end{cases}$$

$$\text{typeg2} = \begin{cases} 1 & \text{if } \text{type_group3} = \text{"Physical/Material"} \\ 0 & \text{otherwise} \end{cases}$$

$$\text{typeg3} = \begin{cases} 1 & \text{if } \text{type_group3} = \text{"Mystical/Supernatural"} \\ 0 & \text{otherwise} \end{cases}$$

Hint: use `ifelse()` or any other functions/methods

Solutions

- We create the requested variables

```
pok <- pok |>
  mutate(
    legend1 = ifelse(legendary == "Yes", 1, 0), legend0 = ifelse(legendary == "No", 1, 0)
  ) |>
  mutate(typeg1 = 1 * (type_group3 == "Elemental/Environmental")) |>
  mutate(typeg2 = 1 * (type_group3 == levels(type_group3)[2])) |>
  mutate(typeg3 = 1 * (type_group3 == levels(type_group3)[3]))
```

- Quick check with `count()` from `{dplyr}`

```
count(pok, legendary, legend0, legend1) |> as.data.frame()
```

	legendary	legend0	legend1	n
1	No	1	0	735
2	Yes	0	1	65

```
count(pok, type_group3, typeg1, typeg2, typeg3) |> as.data.frame()
```

	type_group3	typeg1	typeg2	typeg3	n
1	Elemental/Environmental	1	0	0	334
2	Physical/Material	0	1	0	297
3	Mystical/Supernatural	0	0	1	169

Question 7: Using dummy variables in OLS regression

You have created the dummy variables `legend1`, `legend0`, `typeg1`, `typeg2`, and `typeg3`, which encode information about whether a Pokémon is Legendary and which primary type-group it belongs to.

We now want to explore how these characteristics relate to the `attack` variable using simple and multiple linear regressions.

Using the Pokémon dataset and the dummy variables you created, estimate the following three OLS regression models, each with `attack` as the dependent variable:

$$\text{attack} = \beta_0 + \beta_1 \text{legend1} + \varepsilon \quad (\text{Model 1})$$

$$\text{attack} = \beta_0 + \beta_1 \text{typeg2} + \beta_2 \text{typeg3} + \varepsilon \quad (\text{Model 2})$$

$$\text{attack} = \beta_0 + \beta_1 \text{legend1} + \beta_2 \text{typeg2} + \beta_3 \text{typeg3} + \varepsilon \quad (\text{Model 3})$$

1. For each model, report the regression output and interpret the coefficients.
2. Compare the 3 models with `compare_performance()` from `{performance}`.
3. Refit the 3 models using the factor variables `legendary` and `type_group3`.

Solutions

- We fit the 3 models

```
mod1 <- lm(attack ~ legend1, data = pok)
```

- This specification implies that non-legendsy Pokémon is the reference group

```
mod2 <- lm(attack ~ typeg2 + typeg3, data = pok)
```

- This specification implies that Elemental/Environmental Pokémon is the reference group

```
mod3 <- lm(attack ~ legend1 + typeg2 + typeg3, data = pok)
```

- This specification implies that non-legendsy Pokémon and Elemental/Environmental are the reference groups

Model 1 with `model_parameters()` from `{parameters}`

```
model_parameters(mod1, ci_method = "residual", digits = 1)
```

Parameter	Coefficient	SE	95% CI	t(798)	p
(Intercept)	75.7	1.1	[73.5, 77.9]	67.3	< .001
legend1	41.0	3.9	[33.3, 48.7]	10.4	< .001

$\text{attack} = \beta_0 + \beta_1 \text{legend1} + \varepsilon$
Profile $\mathbb{E}(\text{attack} \cdot)$
Non-legends Pokémons β_0
Legends Pokémons $\beta_0 + \beta_1$

- Intercept: $\hat{\beta}_0 = 75.7$: mean attack for non-legends Pokémons.
- legend1: $\hat{\beta}_1 = 41.0$: difference in mean attack between legends and non-legends Pokémons is on average equal to 41.0 ($p < 0.001$).

Model 2

```
model_parameters(mod2, ci_method = "residual", digits = 1)
```

Parameter	Coefficient	SE	95% CI	t(797)	p
(Intercept)	74.9	1.8	[71.5, 78.4]	42.4	< .001
typeg2	7.1	2.6	[2.0, 12.1]	2.8	0.006
typeg3	6.8	3.1	[0.8, 12.8]	2.2	0.027

$\text{attack} = \beta_0 + \beta_1 \text{typeg2} + \beta_2 \text{typeg3} + \varepsilon$
Profile $\mathbb{E}(\text{attack} \cdot)$
Elemental/Environmental β_0
Physical/Material $\beta_0 + \beta_1$
Mystical/Supernatural $\beta_0 + \beta_2$

- Intercept: $\hat{\beta}_0 = 74.9$: mean attack for Pokémons in the Elemental/Environmental group (reference).
- typeg2: $\hat{\beta}_1 = 7.1$: difference in mean attack between the Physical/Material group and the Elemental/Environmental group is on average equal to 7.1 ($p = 0.006$).
- typeg3: $\hat{\beta}_2 = 6.8$: difference in mean attack between the Mystical/Supernatural group and the Elemental/Environmental group is on average equal to 6.8 ($p = 0.027$).

Model 3

```
model_parameters(mod3, ci_method = "residual", digits = 1)
```

Parameter	Coefficient	SE	95% CI	t(796)	p
(Intercept)	72.4	1.7	[69.1, 75.7]	43.2	< .001
legend1	41.8	4.0	[34.0, 49.7]	10.5	< .001
typeg2	7.6	2.4	[2.9, 12.4]	3.2	0.002
typeg3	1.6	2.9	[-4.1, 7.3]	0.5	0.583

$$\text{attack} = \beta_0 + \beta_1 \text{legend1} + \beta_2 \text{typeg2} + \beta_3 \text{typeg3} + \varepsilon$$

Profile	$\mathbb{E}(\text{attack} \cdot)$
Non-legends, Elemental/Environmental	β_0
Non-legends, Physical/Material	$\beta_0 + \beta_2$
Non-legends, Mystical/Supernatural	$\beta_0 + \beta_3$
Legends, Elemental/Environmental	$\beta_0 + \beta_1$
Legends, Physical/Material	$\beta_0 + \beta_1 + \beta_2$
Legends, Mystical/Supernatural	$\beta_0 + \beta_1 + \beta_3$

- Intercept: $\hat{\beta}_0 = 72.4$: mean attack for non-legends in the Elemental/Environmental group (reference category).
- legend1: $\hat{\beta}_1 = 41.8$: difference in mean attack between the legends and non-legends is on average equal to 41.8 ($p < 0.001$), for all primary type.
- typeg2: $\hat{\beta}_2 = 7.6$: difference in mean attack between the Physical/Material group and the Elemental/Environmental group is on average equal to 7.6 ($p = 0.002$), for all legendary status.
- typeg3: $\hat{\beta}_3 = 1.6$: difference in mean attack between the Mystical/Supernatural group and the Elemental/Environmental group is on average equal to 1.6 ($p = 0.583$), for all legendary status.

- We compare the 3 models

```
compare_performance(
  mod1, mod2, mod3,
  metrics = c("AIC", "BIC", "R2", "R2_adj", "SIGMA", "RMSE")
)
```

Comparison of Model Performance Indices

Name	Model	AIC (weights)	BIC (weights)	R2	R2 (adj.)	RMSE	Sigma
mod1	lm	7741.5 (0.038)	7755.6 (0.811)	0.119	0.118	30.441	30.479
mod2	lm	7836.1 (<.001)	7854.9 (<.001)	0.011	0.009	32.254	32.315
mod3	lm	7735.1 (0.962)	7758.5 (0.189)	0.131	0.127	30.242	30.318

- We perform nested F tests

```
anova(mod1, mod3)
```

```
linearHypothesis(mod3, c("typeg2", "typeg3"))
```

```
Linear hypothesis test:
typeg2 = 0
typeg3 = 0

Model 1: restricted model
Model 2: attack ~ legend1 + typeg2 + typeg3

  Res.Df    RSS Df Sum of Sq   F Pr(>F)
1    798 741307
2    796 731682  2      9624 5.24 0.0055 **
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
anova(mod2, mod3)
```

```
Analysis of Variance Table

Model 1: attack ~ typeg2 + typeg3
Model 2: attack ~ legend1 + typeg2 + typeg3
  Res.Df    RSS Df Sum of Sq   F     Pr(>F)
1    797 832263
2    796 731682  1    100580 109 <0.0000000000000002 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

- We refit the 3 models using the factor variables `legendary` and `type_group3`

```
mod1bis <- lm(attack ~ legendary, data = pok)
model_parameters(mod1bis, digits = 1, include_reference = TRUE)
```

Parameter	Coefficient	SE	95% CI	t(798)	p
(Intercept)	75.7	1.1	[73.5, 77.9]	67.3	< .001
legendary [No]	0.0				
legendary [Yes]	41.0	3.9	[33.3, 48.7]	10.4	< .001

```
mod2bis <- lm(attack ~ type_group3, data = pok)
model_parameters(mod2bis, digits = 1, include_reference = TRUE)
```

Parameter	Coefficient	SE	95% CI	t(797)	p
(Intercept)	74.9	1.8	[71.5, 78.4]	42.4	< .001
type group3 [Elemental/Environmental]	0.0				
type group3 [Physical/Material]	7.1	2.6	[2.0, 12.1]	2.8	0.006
type group3 [Mystical/Supernatural]	6.8	3.1	[0.8, 12.8]	2.2	0.027

```
mod3bis <- lm(attack ~ legendary + type_group3, data = pok)
model_parameters(mod3bis, digits = 1, include_reference = TRUE)
```

Parameter	Coefficient	SE	95% CI	t(796)	p
(Intercept)	72.4	1.7	[69.1, 75.7]	43.2	< .001
legendary [No]	0.0				
legendary [Yes]	41.8	4.0	[34.0, 49.7]	10.5	< .001
type group3 [Elemental/Environmental]	0.0				
type group3 [Physical/Material]	7.6	2.4	[2.9, 12.4]	3.2	0.002
type group3 [Mystical/Supernatural]	1.6	2.9	[-4.1, 7.3]	0.5	0.583

- With this parametrization, the function `Anova` from `{car}` is very useful to test the significance of every variable in the model

```
Anova(mod3bis, type = 3)
```

Anova Table (Type III tests)

```
Response: attack
          Sum Sq Df F value    Pr(>F)
(Intercept) 1716647  1 1867.55 <0.0000000000000002 ***
legendary   100580  1 109.42 <0.0000000000000002 ***
type_group3  9624   2   5.24      0.0055 **
Residuals   731682 796
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

- You can change the reference group with `relevel()` or with `c()`. For example, if you want Mystical/Supernatural as the reference for primary type (type_group3)

```
lm(attack ~ legendary + relevel(type_group3, ref = 3), data = pok) |>
  model_parameters() |>
  format_table(select = "{estimate} [{ci}]|{p}", digits = 1)
```

	Parameter	Coefficient [CI]	p
1	(Intercept)	74.0 [69.2, 78.8]	<0.001
2	legendar	41.8 [34.0, 49.7]	<0.001
3	relevel(type group3, ref = 3)Elemental/Environmental	-1.6 [-7.3, 4.1]	0.583
4	relevel(type group3, ref = 3)Physical/Material	6.0 [0.2, 11.9]	0.043

```
lm(attack ~ legendary + C(type_group3, base = 3), data = pok) |>
  model_parameters() |>
  format_table(select = "{estimate} [{ci}]|{p}", digits = 1)
```

	Parameter	Coefficient [CI]	p
1	(Intercept)	74.0 [69.2, 78.8]	<0.001
2	legendar	41.8 [34.0, 49.7]	<0.001
3	C(type group3, base = 3) [Elemental/Environmental]	-1.6 [-7.3, 4.1]	0.583
4	C(type group3, base = 3) [Physical/Material]	6.0 [0.2, 11.9]	0.043

- We can also use the constraint $\sum \alpha_i = 0$ (course notation).

```
lm(attack ~ legendary + C(type_group3, contr = sum), data = pok) |>
  model_parameters() |>
  format_table(select = "{estimate} [{ci}]|{p}", digits = 1)
```

	Parameter	Coefficient [CI]	p
1	(Intercept)	75.5 [73.2, 77.8]	<0.001
2	legendar	41.8 [34.0, 49.7]	<0.001
3	C(type group3, contr = sum) [1]	-3.1 [-6.0, -0.2]	0.038
4	C(type group3, contr = sum) [2]	4.6 [1.6, 7.5]	0.003

- See the course for the interpretation of coefficients. Some statisticians dislike this parameterization.

Question 8: Testing equality of coefficients

Consider OLS model (mod3):

$$\widehat{\text{attack}} = 72.435 + 41.84 \times \text{legend1} + 7.623 \times \text{typeg2} + 1.595 \times \text{typeg3}$$

Parameter	Coefficient	SE	95% CI	t(796)	p
<hr/>					
(Intercept)	72.435	1.676	[69.145, 75.725]	43.215	< .001
legend1	41.840	4.000	[33.988, 49.691]	10.460	< .001
typeg2	7.623	2.419	[2.876, 12.371]	3.152	0.002
typeg3	1.595	2.904	[-4.106, 7.296]	0.549	0.583

Perform the following test

$$H_0 : \beta_2 = \beta_3 \quad \text{versus} \quad H_1 : \beta_2 \neq \beta_3.$$

Solutions

- We can perform the test with `linearHypothesis()` from `{car}`

```
linearHypothesis(mod3, c("typeg2 = typeg3")) |> as_tibble()
```

```
# A tibble: 2 x 6
  Res.Df    RSS    Df `Sum of Sq`   F `Pr(>F)`
  <dbl>  <dbl> <dbl>      <dbl> <dbl>     <dbl>
1    797 735465.    NA        NA  NA     NA
2    796 731682.    1    3783.  4.12  0.0428
```

- Conclusion: β_2 is significantly different from β_3 .
There is difference in power between Physical/Material and Mystical/Supernatural.

- We can perform the test with `glht()` from `{multcomp}`. We also get the estimated difference and CI

```
glht(mod3, linfct = c("typeg2 - typeg3 = 0")) |>
  model_parameters(digits = 3, verbose = FALSE)
```

```
# Fixed Effects
```

Parameter	Coefficient	SE	95% CI	t(796)	p
typeg2 - typeg3 == 0	6.029	2.972	[0.195, 11.862]	2.029	0.043

- What if we use factor variables in the model ?

```
model_parameters(mod3bis, digits = 1, verbose = FALSE)
```

Parameter	Coefficient	SE	95% CI	t(796)	p
(Intercept)	72.4	1.7	[69.1, 75.7]	43.2	< .001
legendary [Yes]	41.8	4.0	[34.0, 49.7]	10.5	< .001
type_group3 [Physical/Material]	7.6	2.4	[2.9, 12.4]	3.2	0.002
type_group3 [Mystical/Supernatural]	1.6	2.9	[-4.1, 7.3]	0.5	0.583

- Name of the model coefficients with `find_parameters()` from `{insight}`

```
name_param <- find_parameters(mod3bis, flatten = TRUE)
name_param
```

```
[1] "(Intercept)"                  "legendaryYes"
[3] "type_group3Physical/Material" "type_group3Mystical/Supernatural"
```

- Test with `linearHypothesis()`

```
linearHypothesis(mod3bis, glue("{name_param[3]} = {name_param[4]}"))
as.data.frame()
```

Res.Df	RSS	Df	Sum of Sq	F	Pr(>F)
1	797	735465	NA	NA	NA
2	796	731682	1	3783	4.1155 0.042824

- Test with `glht()`

```
C <- rbind("beta2 - beta3" = c(0, 0, 1, -1))
C
```

```
[,1] [,2] [,3] [,4]
beta2 - beta3 0 0 1 -1
```

```
glht(mod3bis, linfct = C) |>
  model_parameters(digits = 3, verbose = FALSE)
```

```
# Fixed Effects
```

Parameter	Coefficient	SE	95% CI	t(796)	p
beta2 - beta3 == 0	6.029	2.972	[0.195, 11.862]	2.029	0.043

Question 9: Predicting mean attack for all category combinations

Consider the following regression model estimated using the Pokémon dataset (`mod3bis`):

$$\text{attack} = \beta_0 + \beta_1 \text{legendary}_{\text{Yes}} + \beta_2 \text{type_group3}_{\text{Physical/Material}} + \beta_3 \text{type_group3}_{\text{Mystical/Supernatural}} + \varepsilon$$

Parameter	Coefficient	SE	95% CI	t(796)	p
(Intercept)	72.4	1.7	[69.1, 75.7]	43.2	< .001
legendary [Yes]	41.8	4.0	[34.0, 49.7]	10.5	< .001
type group3 [Physical/Material]	7.6	2.4	[2.9, 12.4]	3.2	0.002
type group3 [Mystical/Supernatural]	1.6	2.9	[-4.1, 7.3]	0.5	0.583

We want to compute the predicted mean Attack for every combination of these two variables.

1. Create a prediction grid containing all $2 \times 3 = 6$ combinations of
 - `legendary` $\in \{\text{Yes}, \text{No}\}$
 - `type_group3` $\in \{\text{Elemental/Environmental, Physical/Material, Mystical/Supernatural}\}$
2. Compute the predicted mean of `attack` for each of the six combinations.
3. Present the results in a table showing:
 - the combination of categories
 - the predicted mean Attack with standard error or 95% confidence interval

Solutions

- Create all $2 \times 3 = 6$ combinations of `legendary` and `type_group3`

```
newdata_grid <- expand_grid(  
  legendary = fct_unique(pok$legendary),  
  type_group3 = fct_unique(pok$type_group3)  
) |>  
  arrange(legendary, type_group3)  
  
newdata_grid
```

```
# A tibble: 6 x 2  
  legendary type_group3  
  <fct>     <fct>  
1 No        Elemental/Environmental  
2 No        Physical/Material  
3 No        Mystical/Supernatural  
4 Yes       Elemental/Environmental  
5 Yes       Physical/Material  
6 Yes       Mystical/Supernatural
```

- Get predictions with standard errors and CI

```

predictions <- predict(
  mod3bis,
  newdata = newdata_grid, se.fit = TRUE, interval = "confidence"
)

pred_means_tidy <- bind_cols(newdata_grid, predictions[["fit"]]) |>
  mutate(across(is.numeric, \((x) style_number(x, digits = 1)))) |>
  mutate("Mean [95% CI]" = glue("{fit} [{lwr}, {upr}]")) |>
  select(Legendary = legendary, "Primary Type" = type_group3, "Mean [95% CI]")

pred_means_tidy

```

```

# A tibble: 6 x 3
  Legendary `Primary Type`      `Mean [95% CI]`
  <fct>     <fct>            <glue>
1 No        Elemental/Environmental 72.4 [69.1, 75.7]
2 No        Physical/Material    80.1 [76.6, 83.5]
3 No        Mystical/Supernatural 74.0 [69.2, 78.8]
4 Yes       Elemental/Environmental 114.3 [106.2, 122.3]
5 Yes       Physical/Material    121.9 [113.7, 130.1]
6 Yes       Mystical/Supernatural 115.9 [108.0, 123.7]

```

- With `estimate_expectation()` from `{modelbased}`

```

estimate_expectation(mod3bis, data = newdata_grid, ci = 0.95) |>
  rename(Legendary = 1, "Primary Type" = type_group3)

means <- estimate_expectation(mod3bis, by = c("type_group3", "legendary"), ci = 0.95) |>
  rename(Legendary = 1, "Primary Type" = type_group3)

means

```

Model-based Predictions								
Primary Type		legendary		Predicted		SE		95% CI
Elemental/Environmental		No		72.43		1.68		[69.14, 75.72]
Physical/Material		No		80.06		1.77		[76.58, 83.53]
Mystical/Supernatural		No		74.03		2.44		[69.23, 78.83]
Elemental/Environmental		Yes		114.27		4.11		[106.21, 122.34]
Physical/Material		Yes		121.90		4.20		[113.66, 130.14]
Mystical/Supernatural		Yes		115.87		4.01		[107.99, 123.75]

Variable predicted: attack
 Predictors modulated: type_group3, legendary

- Contrast analysis can be achieved through `estimate_contrasts()` from `{modelbased}`

```
predictors <- c("legendary", "type_group3")

mod3bis |>
  estimate_contrasts(contrast = predictors, comparison = "pairwise", p_adjust = "bonferroni") |>
  format_table(select = "{estimate} [{ci}]|{p}", digits = 1) |>
  kable(format = "pipe", align = "l")
```

Level1	Level2	Difference [CI]	p
No, Physical/Material	No, Elemental/Environmental	7.6 [2.9, 12.4]	0.025
No, Mystical/Supernatural	No, Elemental/Environmental	1.6 [-4.1, 7.3]	> .999
Yes, Elemental/Environmental	No, Elemental/Environmental	41.8 [34.0, 49.7]	<0.001
Yes, Physical/Material	No, Elemental/Environmental	49.5 [40.2, 58.7]	<0.001
Yes, Mystical/Supernatural	No, Elemental/Environmental	43.4 [34.6, 52.3]	<0.001
No, Mystical/Supernatural	No, Physical/Material	-6.0 [-11.9, -0.2]	0.642
Yes, Elemental/Environmental	No, Physical/Material	34.2 [25.1, 43.3]	<0.001
Yes, Physical/Material	No, Physical/Material	41.8 [34.0, 49.7]	<0.001
Yes, Mystical/Supernatural	No, Physical/Material	35.8 [26.9, 44.7]	<0.001
Yes, Elemental/Environmental	No, Mystical/Supernatural	40.2 [29.8, 50.7]	<0.001
Yes, Physical/Material	No, Mystical/Supernatural	47.9 [37.3, 58.5]	<0.001
Yes, Mystical/Supernatural	No, Mystical/Supernatural	41.8 [34.0, 49.7]	<0.001
Yes, Physical/Material	Yes, Elemental/Environmental	7.6 [2.9, 12.4]	0.025
Yes, Mystical/Supernatural	Yes, Elemental/Environmental	1.6 [-4.1, 7.3]	> .999
Yes, Mystical/Supernatural	Yes, Physical/Material	-6.0 [-11.9, -0.2]	0.642

Question 10. Residual diagnostics


Using `mod3bis` and functions from the file `helper_functions3.R`:

1. Plot residuals vs fitted values and vs each predictor.
2. Plot $\sqrt{|\text{Standardized residuals}|}$ vs fitted values and vs each predictor.
3. Plot studentized residuals vs fitted values and vs each predictor.
4. Plot residuals vs order of observation.
5. Plot a histogram and normal Q-Q plot of the standardized residuals.

Solutions


- Residuals vs fitted

```
resid_vs_fit(mod3bis, which = "res", jitter = TRUE)
```




- Residuals vs predictors

```
resid_vs_var_factor(mod3bis, predictors, which = "res", fill = "red") +
  theme(axis.text.x = element_text(size = 7))
```




- Studentized residuals vs fitted values

```
resid_vs_fit(mod3bis, which = "rstud", jitter = TRUE)
```



- Residuals vs predictors

```
resid_vs_var_factor(mod3bis, predictors, which = "rstud", fill = "red", size = 1.5) +
  theme(axis.text.x = element_text(size = 7))
```




- Residuals vs observation order. Histogram, Density Q–Q plot of standardized residuals

```
p1 <- resid_vs_order(mod3bis)
p2 <- resid_stand_hist(mod3bis)
p3 <- resid_stand_dens(mod3bis)
p4 <- resid_stand_qq(mod3bis)
```

- We combine the 4 plots with the help of `+` from `{patchwork}`

```
p1 + p2 + p3 + p4
```



Session Info

Package	Version
broom	1.0.10
car	3.1-3
collapse	2.1.4
correlation	0.8.8
datawizard	1.3.0
effectsize	1.0.1
GGally	2.4.0
ggfortify	0.4.19
ggpubr	0.6.2
glue	1.8.0
gtsummary	2.4.0
insight	1.4.2
janitor	2.2.1
kableExtra	1.4.0
lmtest	0.9-40
matrixTests	0.2.3.1
modelbased	0.13.0
multcomp	1.4-29
openxlsx	4.2.8
parameters	0.28.2
patchwork	1.3.2
performance	0.15.2
qqplotr	0.0.7
rstatix	0.7.3
scales	1.4.0
see	0.12.0
tidyverse	2.0.0